

Manoj Kumar

GATE AIR - 13

M.Tech in Data Science From IIT Guwahati

Expertise in Machine Learning, Deep Learning, Artificial Intelligence, Probability and Statistics



Telegram: @Manoj_Gate_DSAI



CLICK



Unacademy



Some True/False ?

1. If events A and B are conditionally independent given C, then A and B are also independent of each other. (True/False)?
2. If A and {B, C} are conditionally independent given D, are A and B conditionally independent given D? (True/False)?
3. Maximizing the likelihood of a logistic regression model yields multiple local optima. (True/False)?
4. For a fixed size of the training and test set, increasing the complexity of the model always leads to a reduction of the test error. (True/False)?
5. Linear regression gives the Maximum Likelihood Estimator (MLE) for data from a model where y_i is normally distributed as $N\left(\sum_j w_j x_{ij}, \sigma^2\right)$. (True/False)?

Some True/False ?

6. Logistic Regression uses the logistic function $\sigma(z) = \frac{1}{1+e^{-z}}$, where z is a linear combination of the input features. Will logistic regression always give a linear decision boundary? (True/False)?
7. In machine learning, the bias is always a bigger source of error than the variance. (True/False)?
8. Suppose a dataset is linearly separable. Is a logistic regressor with regularization parameter $\lambda > 0$ guaranteed to separate the data? (True/False)?
9. In ordinary least squares regression, the coefficient estimator $\hat{w} = (X^T X)^{-1} X^T Y$ that minimizes the residual sum of squares is also the maximum likelihood estimator under the assumption that the errors are normally distributed with constant variance. (True/False)?

Question 10

<https://t.me/gatedsai2025>

In a classification problem with two classes C_1 and C_2 , the posterior probability for class C_1 is defined as:

$$P(C_1|X) = \frac{1}{1+\exp(-a)} = \sigma(a)$$

What is the value of a ?

Options:

1. $a = \ln \left(\frac{P(X|C_1)P(C_1)}{P(X|C_2)P(C_2)} \right)$
2. $a = \ln \left(\frac{P(X|C_2)P(C_2)}{P(X|C_1)P(C_1)} \right)$
3. $a = \ln \left(\frac{P(C_1|X)}{P(X|C_2)P(C_2)} \right)$
4. $a = \ln \left(\frac{P(X|C_2)}{P(X|C_1)} \right)$

<https://t.me/gatedsai2025>

Question 11

<https://t.me/gatedsai2025>

In logistic regression, if the predicted logit value is 0, what is the corresponding transformed probability?

- (A) 0
- (B) 1
- (C) 0.5
- (D) 0.05

Question 12

Consider a naive Bayes classifier with three Boolean input variables, X_1 , X_2 , and X_3 , and one Boolean output variable, Y . How many parameters must be estimated to train this classifier?

<https://t.me/gatedsai2025>

Question 13

<https://t.me/gatedsai2025>

Discriminative models are used to estimate which of the following probabilities? Assume x denotes the input features and y the output labels.

- (a) $P(y \mid x)$
- (b) $P(y, x)$
- (c) $P(x \mid y)$
- (d) $P(y)$

ML by Manoj Kumar

<https://t.me/gatedsai2025>

Question 14

<https://t.me/gatedsai2025>

The solution to the following L_1 regularized least-squares regression problem

$$\arg \min_{\mathbf{w}} \|\mathbf{Y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_1$$

for $\lambda > 0$ is:

(a) $(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}$

(b) $(\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{Y}$

(c) The objective is unbounded, i.e., the solution is $-\infty$.

(d) None of the above

<https://t.me/gatedsai2025>


Question 15

<https://t.me/gatedsai2025>

Suppose you train a logistic regression classifier and the learned hypothesis function is

$$h_{\theta}(x) = \sigma(\theta_0 + \theta_1 x_1 + \theta_2 x_2),$$

where $\theta_0 = 6$, $\theta_1 = 0$, $\theta_2 = -1$. Which of the following represents the decision boundary for $h_{\theta}(x)$?



Question 16

A machine learning model exhibits poor performance on both the training and test datasets. What might explain this outcome?

- A) The model is intricately fitting nuances in the training data, leading to high variance.
- B) The model lacks sufficient complexity to capture the fundamental patterns in the data.
- C) The model performs perfectly on the training data but fails on the test data.
- D) The model's complexity is excessively high relative to the volume of available training data.

Question 17

Which of the following is true about the Naive Bayes classifier?

- A) The Naive Bayes classifier requires that features must be identically distributed within each class to ensure accurate predictions.
- B) The classifier assumes that features are independent across different classes.
- C) The Naive Bayes classifier assumes that all input features are conditionally independent given the class.
- D) The classifier can only be applied to datasets where input features are continuously and normally distributed.

Question 18

<https://t.me/gatedsai2025>

Suppose you are given the following set of data with three Boolean input variables a , b , and c , and a single Boolean output variable K .

a	b	c	K
1	0	1	1
1	1	1	1
0	1	1	0
1	1	0	0
1	0	0	0
0	0	1	1
0	0	1	1
0	0	0	0

Assume we are using a naive Bayes classifier to predict the value of K from the values of the other variables.

what is $P(K = 1 | a = 1, b = 1, c = 0)$?

Question 19[MSQ]

<https://t.me/gatedsai2025>

Given the binary classification task depicted in the figure with the simple linear logistic regression model:

$$P(y = 1|\vec{x}, \vec{w}) = g(w_0 + w_1x_1 + w_2x_2) = \frac{1}{1 + \exp(-w_0 - w_1x_1 - w_2x_2)}$$

Notice that the training data can be separated with zero training error using a linear separator.


Consider training regularized linear logistic regression models where we try to maximize the following objective function:

$$\sum_{i=1}^n \log(P(y_i|x_i, w_0, w_1, w_2)) - Cw_j^2$$

where C is very large and only one of the parameters w_j is regularized in each case.

How does the number of misclassifications change with regularization of each parameter w_j , when C tends to infinity?

- (a) Regularizing w_0 will decrease the number of misclassifications.
- (b) Regularizing w_1 will increase the number of misclassifications.
- (c) Regularizing w_2 will have no effect on the number of misclassifications.
- (d) Regularizing w_0 will increase the number of misclassifications.



Question 20

<https://t.me/gatedsai2025>

Consider the binary classification task depicted in the figure with the simple linear logistic regression model:


$$P(y = 1|\vec{x}, \vec{w}) = g(w_0 + w_1x_1 + w_2x_2) = \frac{1}{1 + \exp(-w_0 - w_1x_1 - w_2x_2)}$$

The training data can be separated with zero training error using a linear separator. Consider training regularized linear logistic regression models where we try to maximize the following objective function:

$$\sum_{i=1}^n \log(P(y_i|x_i, w_0, w_1, w_2)) - C(|w_1| + |w_2|)$$

As we increase the regularization parameter C , which of the following statements can be true?
(Choose only one):

- (a) First w_1 will become 0, then w_2 .
- (b) First w_2 will become 0, then w_1 .
- (c) w_1 and w_2 will become zero simultaneously.
- (d) None of the weights will become exactly zero, only smaller as C increases.



Question 21

<https://t.me/gatedsai2025>

Consider a regularized linear logistic regression model given by:

$$P(y = 1|\mathbf{x}, \mathbf{w}) = g(w_0 + w_1x_1 + w_2x_2)$$

We aim to maximize the following regularized log-likelihood:

$$\sum_{i=1}^n \log p(y_i|\mathbf{x}_i, w_0, w_1, w_2) - Cw_j^2$$

where j can be 0, 1, or 2.

Based on the labeled training set in Figure 4.1 and the resulting training errors (number of misclassifications) as a function of the regularization parameter C in Figure 4.2, assign the "top," "middle," and "bottom" plots to the correct parameter w_0 , w_1 , or w_2 that was regularized in the plot.

- A. Top: w_0 , Middle: w_1 , Bottom: w_2
- B. Top: w_1 , Middle: w_2 , Bottom: w_0
- C. Top: w_2 , Middle: w_1 , Bottom: w_0
- D. Top: w_1 , Middle: w_0 , Bottom: w_2




Figure 4.1 Labeled training set (reproduced here for clarity)




Figure 4.2. Training errors as a function of regularization penalty

Question 22

<https://t.me/gatedsai2025>

Consider the following logistic regression models for binary classification with a sigmoid function

$$g(z) = \frac{1}{1+e^{-z}}:$$

- **Model 1:** $P(Y = 1|X, w_1, w_2) = g(w_1X_1 + w_2X_2)$
- **Model 2:** $P(Y = 1|X, w_1, w_2) = g(w_0 + w_1X_1 + w_2X_2)$

Given three training examples:

- $x^{(1)} = [1, 1]^T, y^{(1)} = 1$
- $x^{(2)} = [1, 0]^T, y^{(2)} = -1$
- $x^{(3)} = [0, 0]^T$, originally $y^{(3)} = 1$

If the label of the third example is changed to -1 , does it affect the learned weights $w = (w_1, w_2)$ in Model 1 and Model 2?

- A. It affects both Model 1 and Model 2.
- B. It affects neither Model 1 nor Model 2.
- C. It affects only Model 1.
- D. It affects only Model 2.

<https://t.me/gatedsai2025>

Question 23

<https://t.me/gatedsai2025>

Consider a logistic regression model where the hypothesis $h_{\theta}(x)$ is given by:

$$h_{\theta}(x) = \sigma(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$$

where $\sigma(z)$ is the sigmoid function defined as $\sigma(z) = \frac{1}{1+e^{-z}}$.

The model includes quadratic transformations of the features x_1^2 and x_2^2 , and a cross-term $x_1 x_2$.

The coefficients for these features are set as follows: $\theta_3 = 2$, $\theta_4 = 2$, and $\theta_5 = 0$. Given these settings, what shape will the decision boundary most likely resemble?

- A) A straight line
- B) A circle or ellipse
- C) A parabola
- D) A hyperbola

<https://t.me/gatedsai2025>

Question 24

<https://t.me/gatedsai2025>

Consider a Gaussian Naive Bayes classifier for a dataset with a single attribute x and two classes 0 and 1. We classify a point \mathbf{x} as class 1 if:

$$P(y = 1|x) \geq P(y = 0|x)$$

For this Gaussian Naive Bayes classifier, the parameters for the two Gaussian distributions are:

$$x|y = 0 \sim N(0, 1/4)$$

$$x|y = 1 \sim N(0, 1)$$

$$P(y = 1) = 0.5$$

Here, $N(\mu, \sigma^2)$ represents a normal distribution with mean μ and variance σ^2 .

Determine the decision boundary for classifying a point \mathbf{x} as class 1.

(a) $x^2 \leq \frac{2 \ln 2}{3}$

(b) $x \leq \sqrt{\frac{2 \ln 2}{3}}$

(c) $x^2 > \frac{2 \ln 2}{3}$

(d) $x \geq \sqrt{\frac{2 \ln 2}{3}}$

<https://t.me/gatedsai2025>

Question 25

<https://t.me/gatedsai2025>

Consider the binary classification problem where class label $Y \in \{0, 1\}$ and each training example X has 2 binary attributes $X_1, X_2 \in \{0, 1\}$. The class priors are $P(Y = 0) = P(Y = 1) = 0.5$, and the conditional probabilities are given as follows:

$P(X_1 Y)$		$X_1 = 0$	$X_1 = 1$
	$Y = 0$	0.7	0.3
	$Y = 1$	0.2	0.8

$P(X_2 Y)$		$X_2 = 0$	$X_2 = 1$
	$Y = 0$	0.9	0.1
	$Y = 1$	0.5	0.5

The expected error rate is the probability that a classifier provides an incorrect prediction for an observation. If Y is the true label, let $\hat{Y}(X_1, X_2)$ be the predicted class label. The expected error rate is:

$$P_D \left(Y \neq \hat{Y}(X_1, X_2) \right) = \sum_{X_1=0}^1 \sum_{X_2=0}^1 P_D \left(X_1, X_2, Y = 1 - \hat{Y}(X_1, X_2) \right)$$

Compute the expected error rate of this naive Bayes classifier which predicts Y given both of the attributes X_1, X_2 . Assume that the classifier is learned with infinite training data.

Question 26

<https://t.me/gatedsai2025>

You are given a dataset with N records in which the i -th record has a real-valued input attribute x_i and a real-valued output y_i , which is generated from a Gaussian distribution with mean $\sin(wx_i)$ and variance 1.

$$P(y_i|w, x_i) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(y_i - \sin(wx_i))^2}{2}\right)$$

We have one unknown parameter w in the above model and we want to learn the maximum likelihood estimate of it from the data.

If you compute the maximum likelihood estimate of the parameter w , which of the following equations is satisfied?

(a) $\sum_i \cos^2(wx_i) = \sum_i y_i \sin(x_i)$

(b) $\sum_i \cos^2(wx_i) = \sum_i y_i \sin(2wx_i)$

(c) $\sum_i x_i \sin(wx_i) \cos(wx_i) = \sum_i y_i \cos(wx_i)$

(d) $\sum_i x_i \cos(x_i) = \sum_i y_i \cos(wx_i)$

Question 27

<https://t.me/gatedsai2025>

Let $F(x) = w_0 + \sum_{j=1}^d w_j x_j$ and $L(y^i F(x^i)) = \frac{1}{1+\exp(y^i F(x^i))}$ be the cost function to be minimized. Here, d represents the number of features, x_j represents the j -th feature of the input x , y is the target variable, and i represents the index of the training sample.

Suppose you use gradient descent to obtain the optimal parameters w_0 and w_j . Give the update rules for these parameters.

Which of the following is the correct update rule for the parameter w_k ?

- (a) $w_k^{(t+1)} = w_k^t - \eta \sum_i x_k^i \left(\frac{\exp(y^i F(x^i))}{(1+\exp(y^i F(x^i)))^2} \right)$
- (b) $w_k^{(t+1)} = w_k^t + \eta \sum_i y^i x_k^i \left(\frac{\exp(y^i F(x^i))}{(1+\exp(y^i F(x^i)))^2} \right)$
- (c) $w_k^{(t+1)} = w_k^t - \eta \sum_i y^i x_k^i \left(\frac{\exp(y^i F(x^i))}{(1+\exp(y^i F(x^i)))^2} \right)$
- (d) $w_k^{(t+1)} = w_k^t + \eta \sum_i x_k^i \left(\frac{\exp(y^i F(x^i))}{(1+\exp(y^i F(x^i)))^2} \right)$